

Building a Movie Database System

1. Before the Change

Query and Execution Plan Without Optimizations

Explain analyze `SELECT "Title", "ReleaseYear", "Genre"`
`FROM movies_partitioned`
`WHERE "Genre" = 'Action' AND "ReleaseYear" > 2015;`

Execution Plan Details:

- Description of the optimizer plan:
The query likely triggers a **full table scan** since no secondary indexes or materialized views are defined. Every row in the `Movies` table is scanned to find matching records.
- Identified Performance Issues:
 - **High resource usage** due to scanning the entire table.
 - **Slow query execution time**, especially with a large dataset.
 - Bottlenecks in data retrieval for frequently queried columns (e.g., `genre`, `release_year`).

Performance Metrics Before Optimizations

- **Execution Time:** 5.066 ms for a dataset of 20000 rows.
- **Resource Usage:** High CPU and I/O activity

Here is a suggested structure for your **Movie Database System Report** based on the submission instructions. Each section addresses the specific requirements mentioned:

2. After the Change

Query and Execution Plan With Optimizations

- **Changes Implemented:**
 - **Secondary Indexes:** Created indexes on frequently filtered columns:
`CREATE INDEX idx_movies_genre ON movies_partitioned ("Genre");`
`CREATE INDEX idx_movies_director ON movies_partitioned ("Director");`
 - **Materialized View:** Precomputed and stored commonly queried data:
`CREATE MATERIALIZED VIEW Top_10_Movies_By_Genre AS`

```
SELECT "Genre", "Title", "Id", "Director", "ReleaseYear", "WatchCount",  
"Rating"  
FROM movies_partitioned  
WHERE "Genre" IS NOT NULL  
ORDER BY "WatchCount" DESC  
LIMIT 10;
```

- **Query Example After Changes:**

```
Explain Analyze SELECT "Title", "Genre", "Director"  
FROM movies_partitioned  
WHERE "Genre" = 'Horror';
```

Execution Plan Details After Optimizations

- **Description of Optimizer Plan:**
 - Use of **indexes** for faster filtering by `genre` and `director`.
 - **Materialized view** directly stores pre-filtered results, avoiding full scans of the base table.
- **Performance Improvements:**
 - Reduced data retrieval time due to precomputed data in the materialized view.
 - Query now uses index lookups instead of full table scans.

Performance Metrics After Optimizations

- **Execution Time:** 1.353 ms "Bitmap Heap Scan on movies_partitioned (cost=36.39..282.26 rows=3110 width=32) (actual time=0.373..1.210 rows=3110 loops=1)"
 - **Resource Usage:** Significant reduction in CPU and I/O activity.
-

3. For Partitioning and API

Partitioning and Performance Improvements

Partitioning Approach:

- Partitioned the `Movies` table by `release_year` for improved performance on year-based queries

Query Execution Plan:

- Queries now target specific partitions, significantly reducing the data scanned.

API Documentation for Movies

This section provides details about the **Movies API** endpoints, including their purpose, methods, request parameters, and example responses.

1. GET /api/Movies

- **Description:** Retrieves a list of all movies in the database.
- **Method:** GET
- **Response Example:**

```
[
  {
    "id": 1,
    "title": "Inception",
    "genre": "Sci-Fi",
    "director": "Christopher Nolan",
    "releaseYear": 2010,
    "watchCount": 1200000,
    "rating": 9.0
  },
  {
    "id": 2,
    "title": "The Dark Knight",
    "genre": "Action",
    "director": "Christopher Nolan",
    "releaseYear": 2018,
    "watchCount": 2000000,
    "rating": 9.1
  },
]
```

2. POST /api/Movies

- **Description:** Adds a new movie to the database.
- **Method:** POST
- **Request Body** (example):

```
{
  "id": 5,
  "title": "The Godfather",
  "genre": "Crime",
  "director": "Francis Ford Coppola",
  "releaseYear": 2020,
  "watchCount": 2500000,
  "rating": 9.2
}
```

3. GET /api/Movies/{id}

- **Description:** Retrieves details for a specific movie by its ID.
- **Method:** GET

- **Path Parameter:**
 - {id}: The unique identifier of the movie (integer).
- **Response Example:**

```
{
  "id": 1,
  "title": "Inception",
  "genre": "Sci-Fi",
  "director": "Christopher Nolan",
  "releaseYear": 2010,
  "watchCount": 1200000,
  "rating": 9.0
}
```

4. PUT /api/Movies/{id}

- **Description:** Updates an existing movie's details.
- **Method:** PUT
- **Path Parameter:**
 - {id}: The unique identifier of the movie to update (integer).
- **Request Body** (example):

```
{
  "id": 1,
  "title": "Inception",
  "genre": "Sci-Fi",
  "director": "Christopher Nolan",
  "releaseYear": 2020 (Updated),
  "watchCount": 47859 (Updated),
  "rating": 9.0
}
```

- **Response Example:**

```
{
  "id": 1,
  "title": "Inception",
  "genre": "Sci-Fi",
  "director": "Christopher Nolan",
  "releaseYear": 2020 (Updated),
  "watchCount": 47859,
  "rating": 9.0
}
```

5. DELETE /api/Movies/{id}

- **Description:** Deletes a specific movie from the database.
- **Method:** DELETE
- **Path Parameter:**
 - {id}: The unique identifier of the movie to delete (integer).
- **Response Example:**

```
{
  "message": "Movie with ID 1 was deleted successfully."
}
```

6. GET /api/Movies/search

- **Description:** Searches for movies based on query parameters.
- **Method:** GET
- **Query Parameters:**
 - genre: Filter by genre.
 - Director: Filter by director.
- **Example Request:**

```
GET api/Movies/search?genre=Comedy&director=Director_246' \
```

Response Example:

```
[
  {
    "id": 20662,
    "title": "Movie_1020661",
    "genre": "Comedy",
    "director": "Director_246",
    "releaseYear": 2010,
    "watchCount": 730,
    "rating": 9.5
  },
  {
    "id": 25310,
    "title": "Movie_1025309",
    "genre": "Comedy",
    "director": "Director_246",
    "releaseYear": 2015,
    "watchCount": 60,
    "rating": 5.5
  }
]
```